

Jiuqi Wang

Email: jiuqi@email.virginia.edu Mobile: +1-434-760-7876 GitHub: [LeonardoWjq](https://github.com/LeonardoWjq)
LinkedIn: [jiuqi-wang-671111196](https://www.linkedin.com/in/jiuqi-wang-671111196) Website: leonardowjq.github.io

EDUCATION

University of Virginia

Ph.D. Computer Science

Advisor: Shangtong Zhang

Aug. 2023 - May 2027

Charlottesville, Virginia

University of Alberta

MSc. Computing Science (Thesis)

Co-Supervisors: Martin Müller and Jonathan Schaeffer

Sep. 2021 - Jun. 2023

Edmonton, Alberta

McGill University

BSc. Honours Computer Science

Distinction, First-Class Honours

Sep. 2017 - Jun. 2021

Montréal, Québec

RESEARCH INTERESTS

Reinforcement Learning, In-Context Learning, Agentic Systems

PUBLICATIONS

* indicates equal contribution; † indicates advisor

PEER-REVIEWED

1. [InfRL: Inference-time Reinforcement Learning for Research Idea Optimization](#)

Sikun Guo, Amir Hassan Shariatmadari, **Jiuqi Wang**, Albert Huang, Stefan Bekiranov, Shangtong Zhang[†], Aidong Zhang

Knowledge Discovery and Data Mining (KDD) AI for Sciences Track, 2026

2. [Almost Sure Convergence of Linear Temporal Difference Learning with Arbitrary Features](#)

Jiuqi Wang, Shangtong Zhang[†]

Journal of Machine Learning Research (JMLR), 2026

3. [Almost Sure Convergence of Differential Temporal Difference Learning for Average Reward Markov Decision Processes](#)

Ethan Blaser, **Jiuqi Wang**, Shangtong Zhang[†]

Poster at **Artificial Intelligence and Statistics (AISTATS)**, 2026

4. [Towards Provable Emergence of In-Context Reinforcement Learning](#)

Jiuqi Wang, Rohan Chandra, Shangtong Zhang[†]

Poster at **Neural Information Processing Systems (NeurIPS)**, 2025

5. [Transformers Learn Temporal Difference Methods for In-Context Reinforcement Learning](#)

Jiuqi Wang*, Ethan Blaser*, Hadi Daneshmand, Shangtong Zhang[†]

Poster at **International Conference on Learning Representations (ICLR)**, 2025

Contributed Talk at the **RLC Workshop on Training Agents with Foundation Models**, 2024

Spotlight Award at the **ICML Workshop on In-Context Learning**, 2024

6. [Deep Dive on Checkers Endgame Data](#)

Jiuqi Wang, Martin Müller[†], Jonathan Schaeffer[†]

Oral presentation at **IEEE Conference on Games (CoG)**, 2023

PREPRINT

1. [Predicting Plasticity in Deep Continual Learning: A Theoretical Perspective](#)

Jiuqi Wang, Jayanth Srinivasa, Claire Chen, Shuze Daniel Liu, Ali Payani, Shangtong Zhang[†]

arXiv:2605.09044, 2026

2. [Rethinking the Design of Reinforcement Learning-Based Deep Research Agents](#)

Yi Wan*, **Jiuqi Wang***, Liam Li, Jinsong Liu, Ruihao Zhu, Zheqing Zhu

arXiv:2510.15862, 2026

Work done during internship at *Pokee AI*; Widely adopted [open-source release](#) accompanying the paper.

3. [Experience Replay Addresses Loss of Plasticity in Continual Learning](#)

Jiuqi Wang, Rohan Chandra, Shangdong Zhang[†]

arXiv:2503.20018, 2025

4. [A Survey of In-Context Reinforcement Learning](#)

Amir Moeini, **Jiuqi Wang**, Jacob Beck, Ethan Blaser, Shimon Whiteson, Rohan Chandra, Shangdong Zhang[†]

arXiv:2502.07978, 2025

RESEARCH & PROFESSIONAL EXPERIENCE

Student Researcher (Upcoming)

May 2026 – Nov. 2026

Google DeepMind

London, UK

- Student Researcher in the autonomous agents team hosted by [Martin Klissarov](#) and [Hao Sun](#).

Research Scientist Intern

Jun. 2025 – Dec. 2025

Pokee AI

Remote

- Co-led the [PokeeResearch-7B](#) open-source project end-to-end.
- Leveraged [verl](#) to efficiently RL-fine-tune and experiment with PokeeResearch-7B.
- Developed tool-server infrastructure and tool-calling schema to support web search and content retrieval.
- Built deployment tools (CLI + [Gradio](#)) for model accessibility and evaluation automation.
- Achieved state-of-the-art performance among 7B-class models on a range of modern Q&A benchmarks; 1.6k stars & 800+ forks on GitHub.

Graduate Research Assistant

Aug. 2023 - present

Advised by Prof. Shangdong Zhang

Charlottesville, Virginia

- Developed theoretical foundations for In-Context Reinforcement Learning (ICRL) in Transformers, resulting in papers accepted at ICLR 2025 and NeurIPS 2025.
- Proved almost-sure convergence of TD learning with arbitrary features; journal paper currently under review.

Master's Thesis – Deep Dive on Checkers Endgame Data

May 2022 - Jun. 2023

Co-Supervised by Prof. Martin Müller and Jonathan Schaeffer

Edmonton, Alberta

- Designed neural network and training pipeline for knowledge compression of checkers endgame data with JAX/Haiku.
- Published as a full paper (oral presentation) at CoG, 2023.

Development Intern

May 2021 – Jul. 2021

PTC Inc.

Suzhou, China

- Improved product internationalization by integrating i18next for multi-language support.

Undergraduate Research

May 2020 - Sep. 2020

Supervised by Prof. Hsiu-chin Lin

Montréal, Québec

- Built and evaluated sim-to-real transfer learning for LPV dynamical systems using Gaussian Mixture Models.
- Developed adaptive retraining procedure to improve robustness in learning with noisy data.

TEACHING ASSISTANTSHIPS

- *Reinforcement Learning (Spring 2026, Fall 2024, University of Virginia)*: hosted office hours and graded assignments.
- *Search, Knowledge and Simulations (Spring 2022, University of Alberta)*: prepared assignment starter code and auto-grading scripts; organized and evaluated the final project (game-playing agent tournament).
- *Intro to the Foundations of Computation I (Fall 2021, University of Alberta)*: held lab sessions to review and complement the course material.
- *Intro to Computer Science (Spring 2020, McGill University)*: prepared original weekly quiz questions.

VOLUNTEERING EXPERIENCE

High School Mentor <i>Charlottesville High School</i>	<i>Sep. 2024 - Dec. 2024</i> Charlottesville, Virginia
• Matched with a team in the senior engineering class and advised their final engineering project. • Held weekly meetings with the team to keep up with the progress and assisted troubleshooting.	

PEER REVIEW

NeurIPS (2025-2026), EWRL (2026), ICML (2024–2026), ICLR (2024–2026), AISTATS (2024-2026), IJCAI (2025-2026), AAMAS (2025–2026), ACML (2024), RLC (2024–2026)

HONORS AND AWARDS

Best Reviewer Award	<i>Apr. 2025</i> AISTATS
AI-SCORE Summer School Fellow <i>\$900</i>	<i>May 2024</i> University of Maryland, College Park
IEEE CIS Travel Grant <i>\$500</i>	<i>Jul. 2023</i> IEEE CIS
GSA Academic Travel Grant <i>\$500</i>	<i>Jul. 2023</i> University of Alberta
Science Undergraduate Research Award (SURA) <i>\$7,000</i>	<i>May 2020</i> McGill University
Dean’s Honour List	<i>Aug. 2018</i> McGill University
Merit-Based Undergraduate Scholarships (Condensed) <i>\$6,500</i>	<i>Jul. 2017 – Jul. 2018</i> McGill University

SKILLS

- **Technology:** Python (NumPy, PyTorch, JAX, vLLM, Gradio), Java, Git/GitHub, Docker, Linux, Weights & Biases, Huggingface, L^AT_EX
- **Knowledge:** Statistical Machine Learning, Deep Learning, (Deep) Reinforcement Learning, Heuristic Search, Robotics
- **Language:** English (proficient), Mandarin (native), French (basic)